

The correctness of an algorithm often depends on properties of one or more data structures. For example, the algorithm of module *Reachable* depends on properties of directed graphs. Assuming those properties, once we have found suitable invariants, proving correctness of the algorithm is often a matter of checking a large number of fairly simple details. Writing a proof checked by *TLAPS* can be a cost-effective way of making sure that the algorithm is correct—assuming that those properties are correct.

Writing *TLAPS* checked proofs of properties of data structures is often difficult. When verifying correctness of an algorithm, it may not be a cost-effective way of verifying correctness of those properties of data structures. Correctness of an algorithm rarely depends on subtle new mathematical properties of the data structures. The properties of a data structure that the algorithm relies on are almost always well known or obvious. The incorrectness of the algorithm's correctness proof introduced by assuming a property of a data structure will most likely be the result of an incorrect TLA+ statement of a correct property. Checking the property on small models is an effective way of catching such an error.

And remember that it's easier to get *TLAPS* to prove something if it's true. Even if you intend to prove properties of a data structure, you should use *TLC* to check that those properties are true before you start writing a TLA+ proof.

The proof of correctness in module *ReachableProofs* of the algorithm in module *Reachable* uses four properties of directed graphs: reachability properties named *Reachable0*, *...*, *Reachable3*. This module defines operators that can be used to have *TLC* check the correctness of those properties.

EXTENDS *Integers*, *FiniteSets*, *TLAPS*, *TLC*

CONSTANTS *Nodes*, *SuccSet*

The algorithm of module *Reachable* has constant parameters *Nodes*, representing a set of nodes, and *Succ*, which is the successor function in [*Nodes* → SUBSET *Nodes*] that defines a directed graph with set *Nodes* of nodes. *SuccSet* is a set of possible values of *Succ* on which *TLC* will test the reachability properties.

$P(Succ) \triangleq$ INSTANCE *ReachabilityProofs*

For every definition *Op* in module *ReachabilityProofs*, this statement defines the operator $P(_)!Op$ in the current module, where the definition of $P(s)!Op$ is obtained from the definition of *Op* in *ReachabilityProofs* by substituting *s* for *Succ* (and substituting the constant *Nodes* of the current module for the constant *Nodes* of module *ReachabilityProofs*).

Warning: *TLC* can handle this kind of parametrized instantiation because *ReachabilityProofs* is a constant module. A bug in the current version of *TLC* makes it unable to handle parametrized instantiation of a non-constant module (one with VARIABLE declarations).

Tests can be performed by having *TLC* evaluate the constant *Test* (using the Evaluate Constant Expression of a model's Model Checking Results page), and checking that its value is $\langle \text{TRUE}, \text{TRUE}, \text{TRUE}, \text{TRUE} \rangle$.

$Test \triangleq \langle \forall Succ \in SuccSet : P(Succ)!Reachable0,$
 $\quad \forall Succ \in SuccSet : P(Succ)!Reachable1,$
 $\quad \forall Succ \in SuccSet : P(Succ)!Reachable2,$
 $\quad \forall Succ \in SuccSet : P(Succ)!Reachable3 \rangle$

The first test to perform is to evaluate *Test* with *SuccSet* equal to the set of all possible values of *Succ*, using as large a set *Nodes* as we can. We do this with a model in which *SuccSet* equals the following set *SuccSet1*.

$SuccSet1 \triangleq [Nodes \rightarrow \text{SUBSET } Nodes]$

For 3 nodes, *SuccSet1* has 3^{2^3} (about 6500) elements. Evaluating *Test* on my laptop takes *TLC* a few seconds—mostly start-up time.

For 4 nodes, *SuccSet1* has 4^{16} (about $4 \cdot 10^9$) elements. I expect it will take *TLC* quite a few hours, and perhaps days, to evaluate *Test*.

We'd like to evaluate *Test* for more than 4 nodes, so we do it by using a model that sets *Succ* to *SuccSet2(n)*, which we now define to be a set of *n* randomly chosen values of *Succ*. A randomly chosen value of *Succ* is a function that assigns to each node *n* a randomly chosen subset of *Nodes*. We choose a randomly chosen subset of *Nodes* by using the *RandomElement* operator from the *TLC* module.

In the following definition, *RandomSucc* is given an unused parameter because *TLC* tries to optimize its execution by evaluating a definition of a constant value once when it starts up and using that same value every time it has to evaluate the constant.

```

RandomSucc(x)  $\triangleq$ 
  LET RECURSIVE RS(–, –)
    RS(F, D)  $\triangleq$  IF D = {}
      THEN F
      ELSE LET d  $\triangleq$  CHOOSE d  $\in$  D : TRUE
        S  $\triangleq$  RandomElement(SUBSET Nodes)
        IN RS(F @@ (d := S), D \ {d})
  IN RS(⟨⟩, Nodes)

```

$SuccSet2(n) \triangleq \{RandomSucc(i) : i \in 1 \dots n\}$

With *Succ* set to *SuccSet2(n)*, it takes *TLC* an average of about $3n$ seconds to evaluate *Test* on my laptop for a set of 5 nodes. It seems to me that an error in the definition of *ReachableFrom* or in one of the lemmas in one of the lemmas *Reachability0* – 3 would almost certainly be manifest in a graph with 5 nodes. So, had these lemmas not been proved, I would have performed as many tests as I could on a graph with 5 nodes.

```

\ * Modification History
\ * Last modified Sun Apr 14 15:36:03 PDT 2019 by lamport
\ * Created Fri Apr 05 19:21:52 PDT 2019 by lamport

```